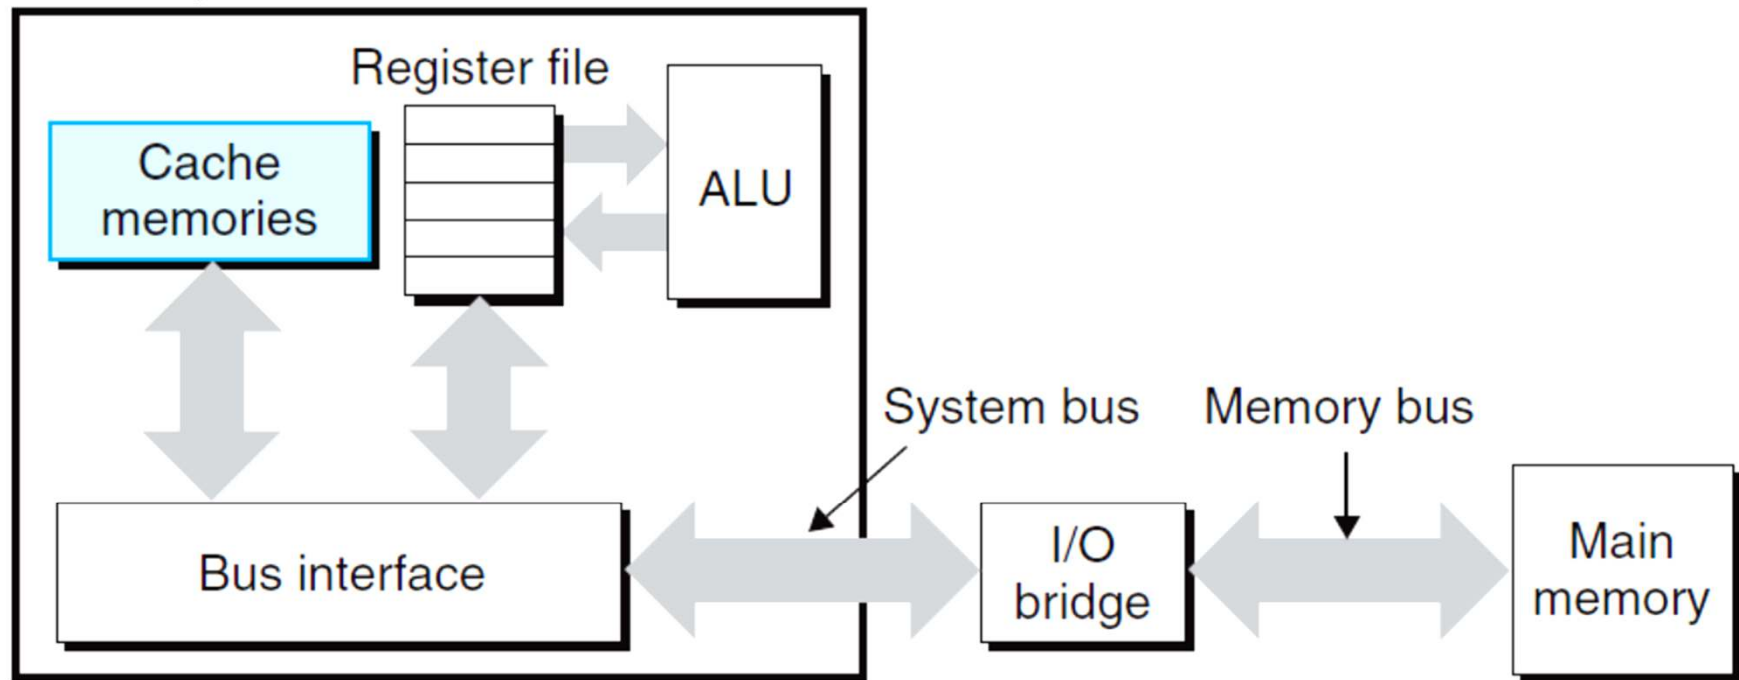




# Typical bus structure for cache memories.

CPU chip

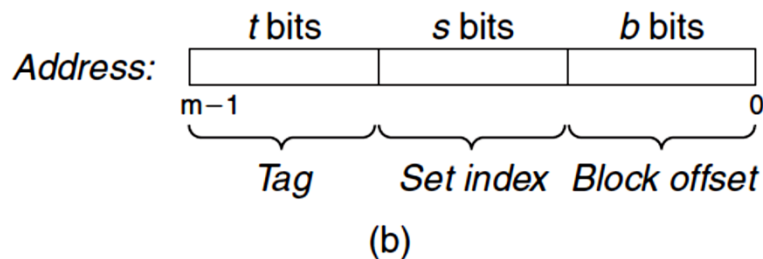
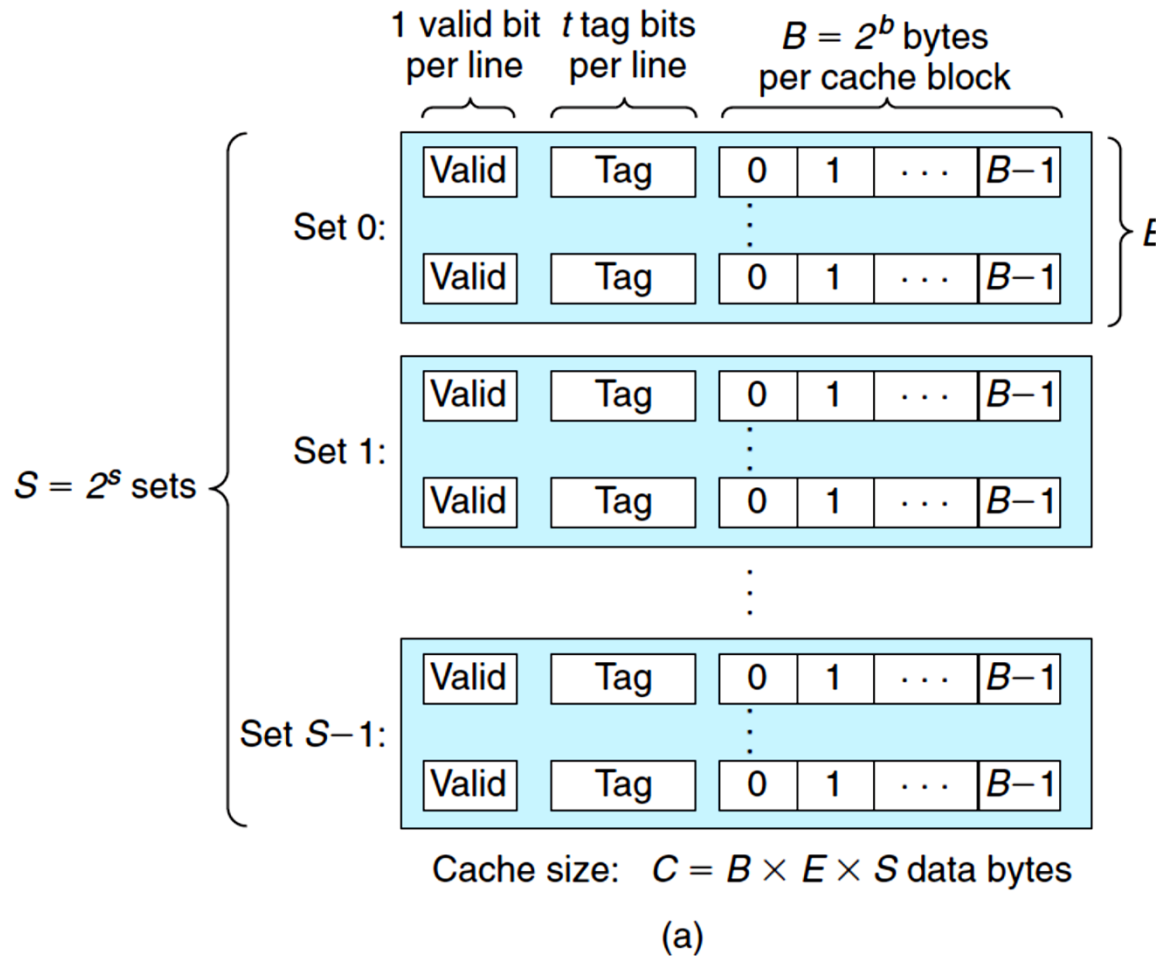


- The memory hierarchies of early computer systems consisted of only three levels:
  - CPU registers
  - Main DRAM memory
  - Disk storage
- However, because of the increasing gap between CPU and main memory, system designers were compelled to insert a small SRAM *cache memory*, called an *L1 cache* (Level 1 cache) between the CPU register file and main memory

## General organization of cache

(a) A cache is an array of sets. Each set contains one or more lines. Each line contains a valid bit, some tag bits, and a block of data.

(b) The cache organization induces a partition of the  $m$  address bits into  $t$  tag bits,  $s$  set index bits, and  $b$  block offset bits.



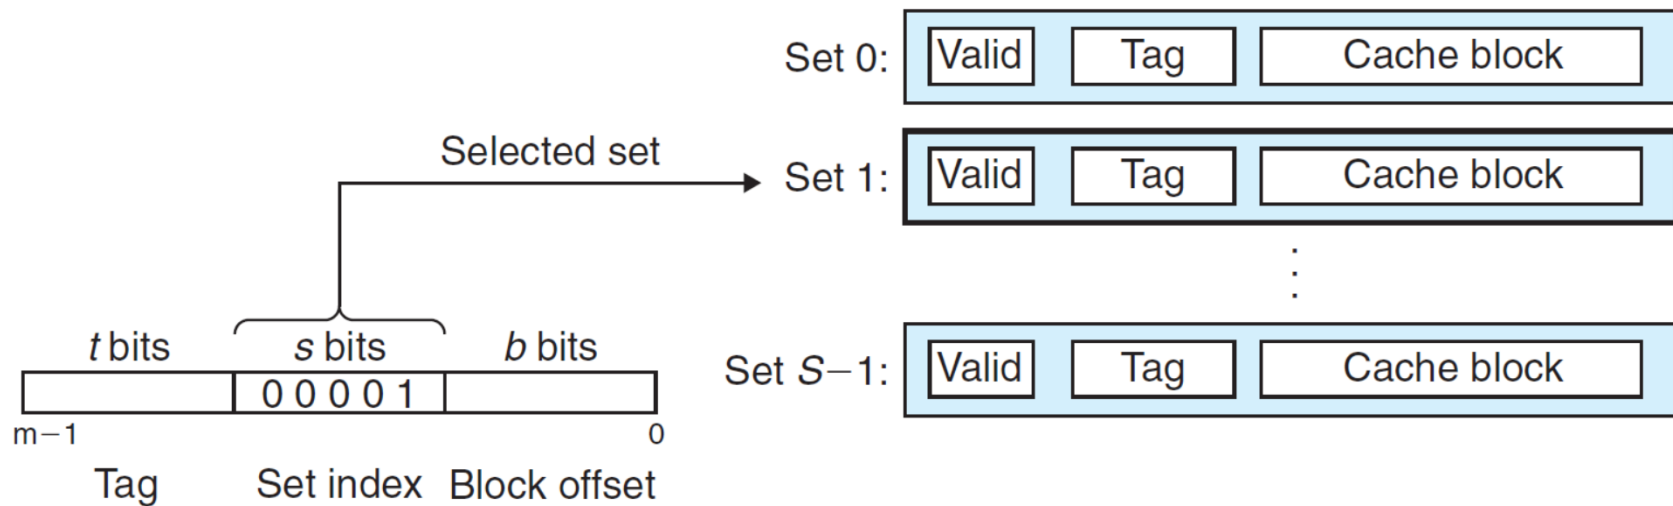
The cache is organized so that it can find the requested word by simply inspecting the bits of the address, similar to a hash table with an extremely simple hash function. Here is how it works:

The parameters  $S$  and  $B$  induce a partitioning of the  $m$  address bits into the three

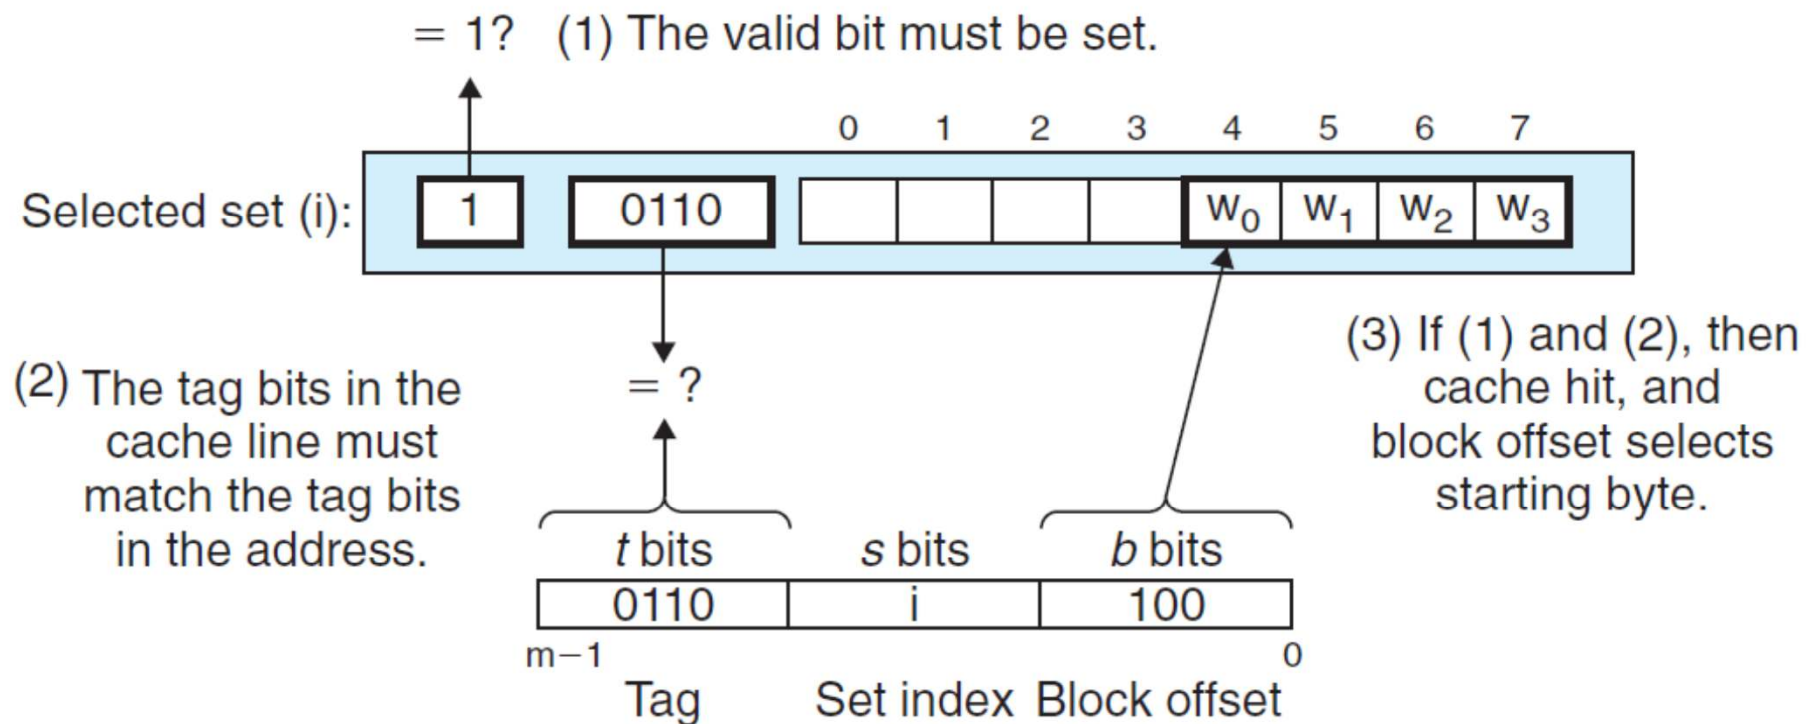
Fundamental parameters	
Parameter	Description
$S = 2^s$	Number of sets
$E$	Number of lines per set
$B = 2^b$	Block size (bytes)
$m = \log_2(M)$	Number of physical (main memory) address bits

- When the CPU is instructed by a load instruction to read a word from address  $A$  of main memory, it sends the address  $A$  to the cache. If the cache is holding a copy of the word at address  $A$ , it sends the word immediately back to the CPU.
- how does the cache know whether it contains a copy of the word at address  $A$ ?

- **Set selection in a direct-mapped cache.**



- Line matching and word selection in a directmapped cache





# A concrete example will help clarify the process

- Suppose we have a direct-mapped cache described by:

$$(S, E, B, m) = (4, 1, 2, 4)$$

- Initially, the cache is empty (i.e., each valid bit is zero):

Set	Valid	Tag	block[0]	block[1]
0	0			
1	0			
2	0			
3	0			

- 1. Read word at address 0. (cache miss)**

Set	Valid	Tag	block[0]	block[1]
0	1	0	m[0]	m[1]
1	0			
2	0			
3	0			

- **Read word at address 1. This is a cache hit**
- **Read word at address 13**

Set	Valid	Tag	block[0]	block[1]
0	1	0	m[0]	m[1]
1	0			
2	1	1	m[12]	m[13]
3	0			

- **Read word at address 8.**

Set	Valid	Tag	block[0]	block[1]
0	1	1	m[8]	m[9]
1	0			
2	1	1	m[12]	m[13]
3	0			

- **Read word at address 0.**

Set	Valid	Tag	block[0]	block[1]
0	1	0	m[0]	m[1]
1	0			
2	1	1	m[12]	m[13]
3	0			

# Conflict Misses in Direct-Mapped Caches

- Conflict misses in direct-mapped caches typically occur when programs access arrays whose sizes are a power of 2
  - For example, consider a function that computes the dot product of two vectors:

```
1  float dotprod(float x[8], float y[8])
2  {
3      float sum = 0.0;
4      int i;
5
6      for (i = 0; i < 8; i++)
7          sum += x[i] * y[i];
8      return sum;
9  }
```

- This function has good spatial locality with respect to x and y, and so we might expect it to enjoy a good number of cache hits. Unfortunately, this is not always true.

- Suppose that  $x$  is loaded into the 32 bytes of contiguous memory starting at address 0, and that  $y$  starts immediately after  $x$  at address 32
- For simplicity, suppose that a block is 16 bytes and that the cache consists of two sets, for a total cache size of 32 bytes.
- We will assume that the variable  $sum$  is actually stored in a CPU register and thus does not require a memory reference

- Given these assumptions, each  $x[i]$  and  $y[i]$  Will map to the identical cache set:

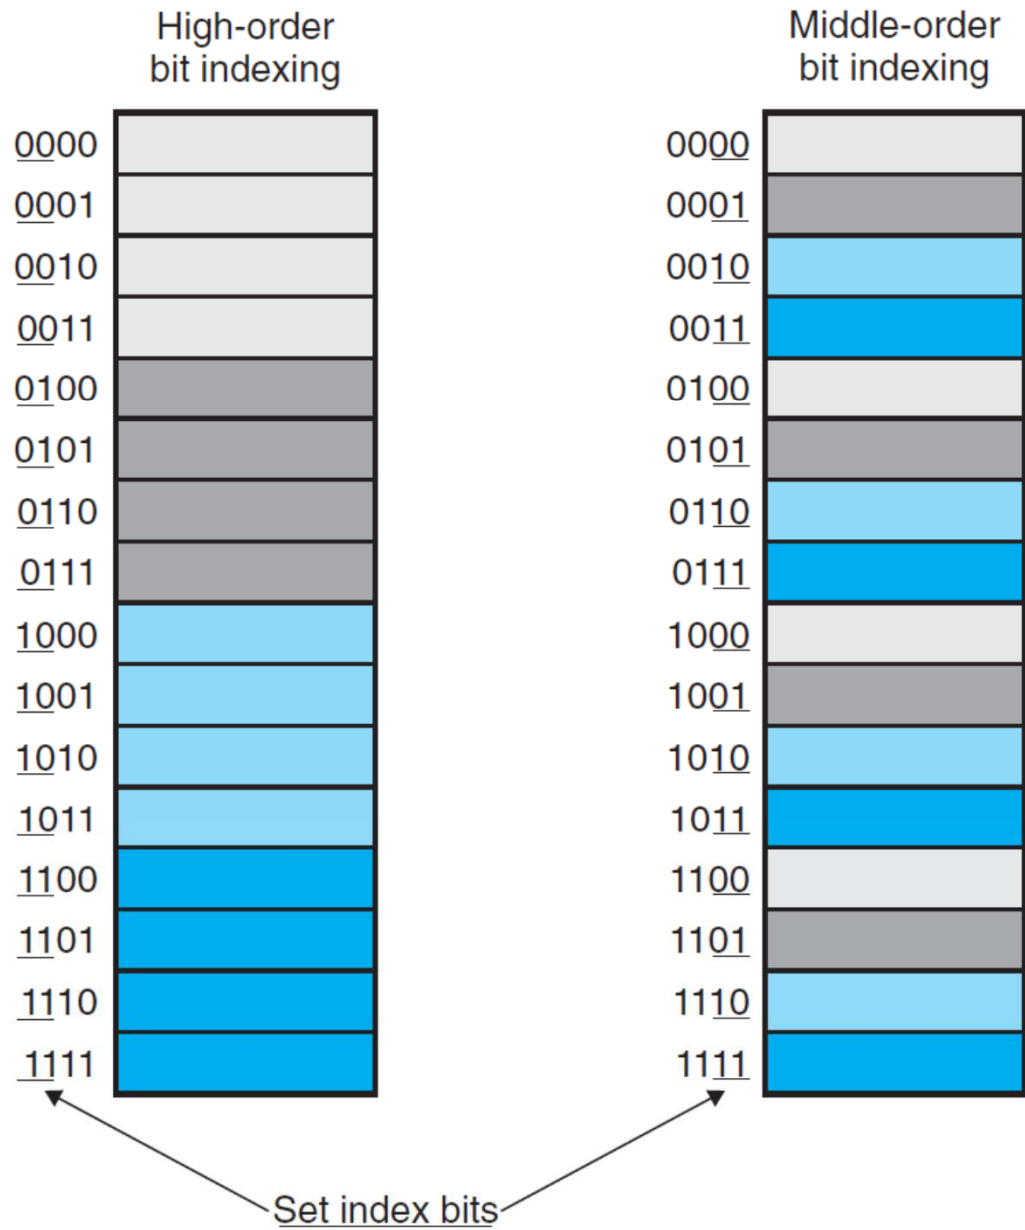

Element	Address	Set index	Element	Address	Set index
$x[0]$	0	0	$y[0]$	32	0
$x[1]$	4	0	$y[1]$	36	0
$x[2]$	8	0	$y[2]$	40	0
$x[3]$	12	0	$y[3]$	44	0
$x[4]$	16	1	$y[4]$	48	1
$x[5]$	20	1	$y[5]$	52	1
$x[6]$	24	1	$y[6]$	56	1
$x[7]$	28	1	$y[7]$	60	1

- One easy solution is to put  $B$  bytes of padding at the end of each array.
- For example, instead of defining  $x$  to be float  $x[8]$ , we define it to be float  $x[12]$ . Assuming  $y$  starts immediately after  $x$  in memory, we have the following mapping of array elements to sets:

Element	Address	Set index	Element	Address	Set index
$x[0]$	0	0	$y[0]$	48	1
$x[1]$	4	0	$y[1]$	52	1
$x[2]$	8	0	$y[2]$	56	1
$x[3]$	12	0	$y[3]$	60	1
$x[4]$	16	1	$y[4]$	64	0
$x[5]$	20	1	$y[5]$	68	0
$x[6]$	24	1	$y[6]$	72	0
$x[7]$	28	1	$y[7]$	76	0

- what fraction of the total references to  $x$  and  $y$  will be hits once we have padded array  $x$ ?





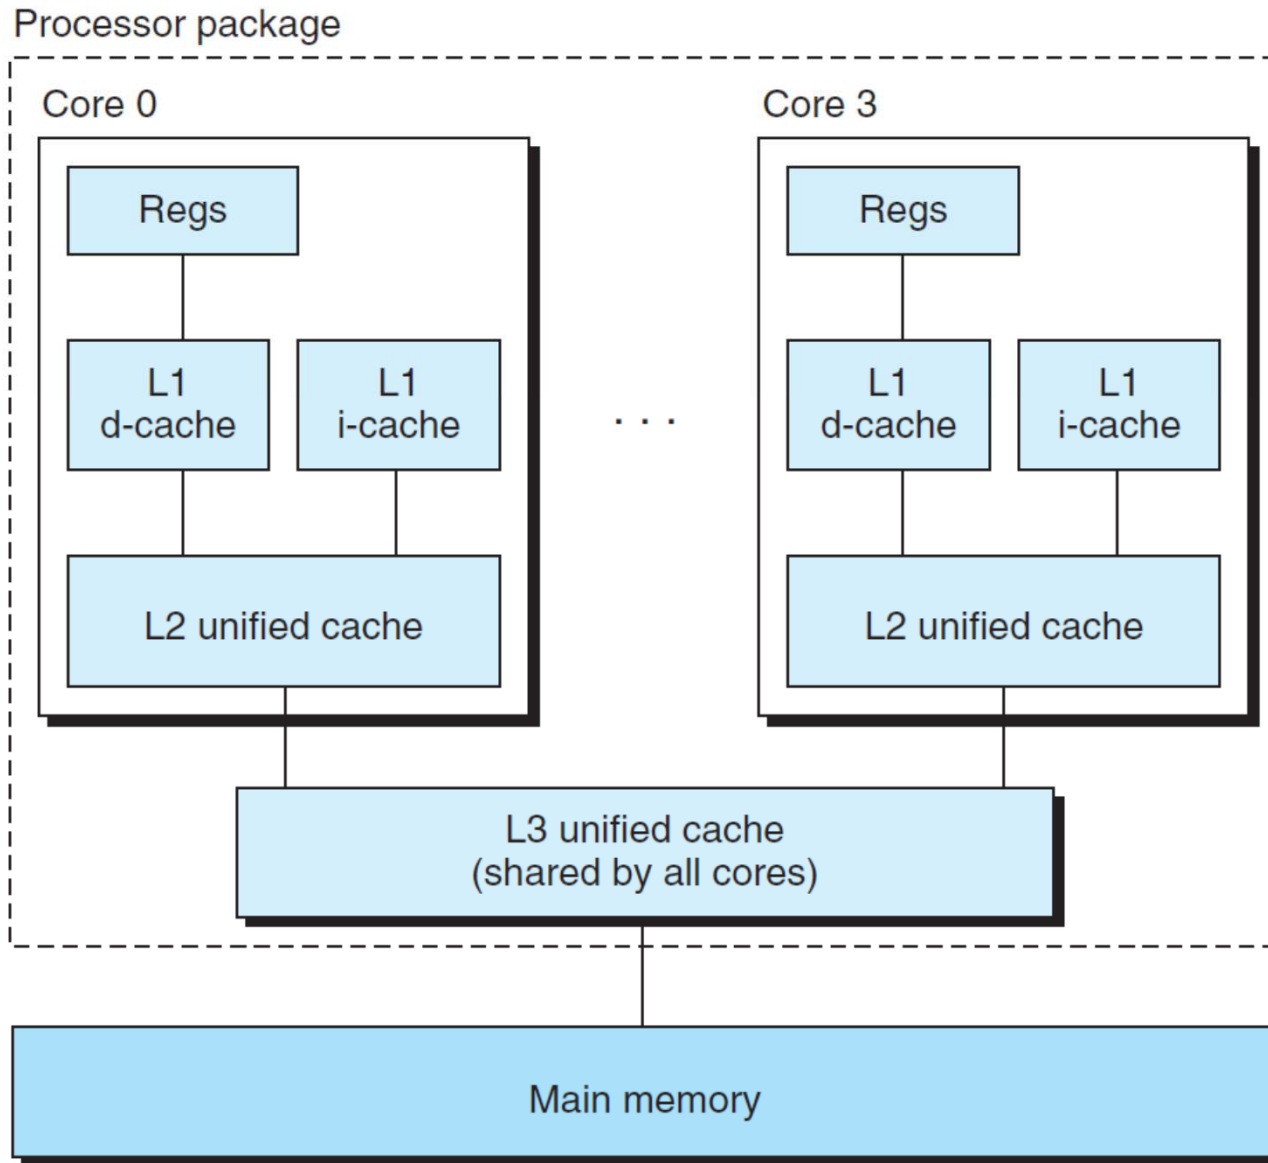
# “cache-friendly” code?

```
for(unsigned int y=0; y<height; ++y)
    { for(unsigned int x=0; x<width; ++x) { ...
      image[y][x] ... } }
```

```
for(unsigned int x=0; x<width; ++x)
    { for(unsigned int y=0; y<height; ++y) { ...
      image[y][x] ... } }
```

# Anatomy of a Real Cache Hierarchy

## Intel Core i7 cache hierarchy.



# Characteristics of the Intel Core i7 cache hierarchy

Cache type	Access time (cycles)	Cache size ( $C$ )	Assoc. ( $E$ )	Block size ( $B$ )	Sets ( $S$ )
L1 i-cache	4	32 KB	8	64 B	64
L1 d-cache	4	32 KB	8	64 B	64
L2 unified cache	11	256 KB	8	64 B	512
L3 unified cache	30–40	8 MB	16	64 B	8192

# Cache-friendly Code

```
1  int sumarrayrows(int a[M][N])
2  {
3      int i, j, sum = 0;
4
5      for (i = 0; i < M; i++)
6          for (j = 0; j < N; j++)
7              sum += a[i][j];
8      return sum;
9  }
```

a[i][j]	j = 0	j = 1	j = 2	j = 3	j = 4	j = 5	j = 6	j = 7
i = 0	1 <b>[m]</b>	2 [h]	3 [h]	4 [h]	5 <b>[m]</b>	6 [h]	7 [h]	8 [h]
i = 1	9 <b>[m]</b>	10 [h]	11 [h]	12 [h]	13 <b>[m]</b>	14 [h]	15 [h]	16 [h]
i = 2	17 <b>[m]</b>	18 [h]	19 [h]	20 [h]	21 <b>[m]</b>	22 [h]	23 [h]	24 [h]
i = 3	25 <b>[m]</b>	26 [h]	27 [h]	28 [h]	29 <b>[m]</b>	30 [h]	31 [h]	32 [h]

```

1  int sumarraycols(int a[M][N])
2  {
3      int i, j, sum = 0;
4
5      for (j = 0; j < N; j++)
6          for (i = 0; i < M; i++)
7              sum += a[i][j];
8      return sum;
9  }

```

$a[i][j]$	$j = 0$	$j = 1$	$j = 2$	$j = 3$	$j = 4$	$j = 5$	$j = 6$	$j = 7$
$i = 0$	1 [m]	5 [m]	9 [m]	13 [m]	17 [m]	21 [m]	25 [m]	29 [m]
$i = 1$	2 [m]	6 [m]	10 [m]	14 [m]	18 [m]	22 [m]	26 [m]	30 [m]
$i = 2$	3 [m]	7 [m]	11 [m]	15 [m]	19 [m]	23 [m]	27 [m]	31 [m]
$i = 3$	4 [m]	8 [m]	12 [m]	16 [m]	20 [m]	24 [m]	28 [m]	32 [m]

- **Solution to Problem 6.20 (page 620)**
- The key to this problem is noticing that the cache can only hold  $1/2$  of the array.
- So the column-wise scan of the second half of the array evicts the lines that
- were loaded during the scan of the first half. For example, reading the first element of `grid[8][0]` evicts the line that was loaded when we read elements from
- `grid[0][0]`. This line also contained `grid[0][1]`. So when we begin scanning the
- next column, the reference to the first element of `grid[0][1]` misses.
- A. What is the total number of read accesses? 512 reads.
- B. What is the total number of read accesses that miss in the cache? 256 misses.
- C. What is the miss rate?  $256/512 = 50\%$ .
- D. What would the miss rate be if the cache were twice as big? If the cache were
- twice as big, it could hold the entire grid array. The only misses would be
- the initial cold misses, and the miss rate would be  $1/4 = 25\%$ .

- You should also assume the following:
  - `sizeof(int) == 4`.
  - grid begins at memory address 0.
  - The cache is initially empty.
  - The only memory accesses are to the entries of the array grid. Variables `i`, `j`, `total_x`, and `total_y` are stored in registers.



- determine the cache performance of the following code:

```
1      for (i = 0; i < 16; i++){  
2          for (j = 0; j < 16; j++) {  
3              total_x += grid[j][i].x;  
4              total_y += grid[j][i].y;  
5          }  
6  }
```

- A. What is the total number of reads?
- B. What is the total number of reads that miss in the cache?
- C. What is the miss rate?
- D. What would the miss rate be if the cache were twice as big?

Considere el problema de la multiplicación de un par de matrices de  $n \times n$  :  $C=A.B$ . Por ejemplo si  $n = 2$ , entonces

- Donde:

$$\begin{bmatrix} c_{11} & c_{12} \\ c_{21} & c_{22} \end{bmatrix} = \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix} \begin{bmatrix} b_{11} & b_{12} \\ b_{21} & b_{22} \end{bmatrix}$$

$$c_{11} = a_{11}b_{11} + a_{12}b_{21}$$

$$c_{12} = a_{11}b_{12} + a_{12}b_{22}$$

$$c_{21} = a_{21}b_{11} + a_{22}b_{21}$$

$$c_{22} = a_{21}b_{12} + a_{22}b_{22}$$